

A Dual-Branch Graph-Temporal Neural Network for Predicting α -Synuclein Pathology Spread Across the Brain Connectome

S. P. Jeni Flemin
Independent Researcher
Puttalam, Sri Lanka
jeniflemin1589@gmail.com

Vigneswaran Vithiyasahar
Faculty of Computing
SLIIT Northern University
Jaffna, Sri Lanka
vithiyasahar.v@slit.lk

Kanagasabai Thiruthanigesan
Faculty of Computing
SLIIT Northern University
Jaffna, Sri Lanka
thiruthanigesan.k@northernuni.lk

Abstract—The prion-like spread of misfolded α -synuclein (α -syn) through the brain connectome is a hallmark of Parkinson’s disease pathogenesis, yet existing network diffusion models cannot capture the nonlinear, region-specific dynamics governing this process. We propose a dual-branch graph-temporal neural network that jointly models spatial propagation patterns via Graph Attention Networks (GATv2) and temporal aggregation kinetics via Gated Recurrent Units (GRU). Evaluated on the Henderson et al. (2019) mouse connectome dataset comprising 426 brain regions across multiple post-injection time points, our optimized model achieves ROC-AUC of 0.780, PR-AUC of 0.821, and accuracy of 74.4%, representing an 8.0% improvement in ROC-AUC over the baseline configuration. A systematic ablation study across seven model variants demonstrates that the spatial-temporal fusion produces a 12.8% ROC-AUC gain over single-branch architectures, confirming that both connectome topology and temporal kinetics are essential for accurate pathology prediction. Attention analysis reveals biologically plausible hub regions—perirhinal cortex, mammillary body, and cingulate cortex—consistent with known α -syn vulnerability patterns.

Index Terms—graph neural networks, α -synuclein, brain connectome, pathology prediction, Parkinson’s disease, GATv2, temporal modeling

I. INTRODUCTION

Parkinson’s disease (PD) is the second most prevalent neurodegenerative disorder, affecting over 10 million people worldwide. A central pathological hallmark is the accumulation and spread of misfolded α -synuclein (α -syn) protein, which forms Lewy body inclusions that propagate through neuronal circuits in a stereotypical pattern first characterized by Braak et al. [1]. Evidence from prion biology [2] and pre-formed fibril (PFF) injection models [3] has established that α -syn pathology spreads via trans-synaptic transmission along axonal pathways [4], suggesting that the structural connectome—the wiring diagram of neural connections—plays a critical role in governing disease progression.

The Network Diffusion Model (NDM) proposed by Raj et al. [5] was the first to formalize this hypothesis, demonstrating that a linear diffusion process on the brain’s structural connectome could predict regional atrophy patterns in dementia. Subsequent work extended this framework to α -syn spread in

mouse models, showing that network connectivity outperforms spatial proximity in predicting pathology distribution [6]. Henderson et al. [7] further validated this approach using a comprehensive PFF injection experiment across 426 mouse brain regions, providing the first large-scale benchmark for connectome-based pathology prediction.

Despite these advances, existing approaches suffer from two key limitations. First, linear diffusion models [8] assume homogeneous spreading rates and cannot capture the nonlinear, region-specific vulnerability patterns observed in vivo [9]. Second, purely spatial models ignore the temporal dynamics of protein aggregation—the kinetic processes by which soluble α -syn converts to pathological forms at different rates across regions.

Graph Neural Networks (GNNs) [10] have emerged as powerful tools for learning on graph-structured data, with Graph Attention Networks (GATs) [11] and their improved variant GATv2 [12] enabling adaptive, attention-weighted message passing. Recent work has demonstrated the utility of GNNs for brain connectivity analysis [13], [14], and multi-modal graph learning has shown promise for disease prediction [15]. However, no prior work has combined graph-based spatial modeling with temporal sequence modeling for connectome-based pathology spread prediction.

In this paper, we propose a **dual-branch graph-temporal neural network** that addresses both limitations through: (1) a *spatial branch* using GATv2 to learn attention-weighted propagation patterns over the brain connectome, capturing region-specific vulnerability via learned attention coefficients; and (2) a *temporal branch* using GRU [16] to model the kinetic trajectory of pathology accumulation at each region. The two branches are fused to produce a joint spatial-temporal representation for per-region pathology prediction. Our contributions are:

- A novel dual-branch architecture combining GATv2-based spatial message passing with GRU-based temporal modeling for brain pathology prediction.
- A comprehensive ablation study across seven model variants demonstrating the necessity of spatial-temporal

fusion.

- Systematic hyperparameter optimization achieving state-of-the-art performance on the Henderson et al. [7] benchmark.
- Interpretable attention analysis revealing biologically plausible vulnerability hubs.

II. RELATED WORK

A. Network Models of Pathology Spread

The concept that neurodegenerative pathology spreads along structural brain networks has gained substantial support from both experimental and computational evidence [8]. The NDM [5] models pathology transmission as a graph diffusion process governed by the Laplacian of the structural connectivity matrix. While validated across multiple diseases [9], the NDM assumes linear, homogeneous dynamics and lacks the capacity to learn nonlinear region-specific effects. The Allen Mouse Brain Connectivity Atlas [17] has provided the structural connectome data underpinning these models. Henderson et al. [7] created the primary benchmark for α -syn spread prediction by quantifying pathology across 426 regions at multiple time points following PFF injection, while Mezas et al. [6] demonstrated the superiority of connectivity-based over proximity-based predictions.

B. Graph Neural Networks for Brain Analysis

GNNs have been increasingly applied to neuroimaging data. Graph Convolutional Networks (GCNs) [10] perform spectral-domain convolutions on graphs, while GATs [11] introduce attention mechanisms for adaptive neighbor aggregation. GATv2 [12] addresses the limited expressiveness of the original GAT by computing dynamic attention that depends on both source and target nodes. BrainGNN [13] demonstrated interpretable graph neural networks specifically for fMRI analysis, and recent reviews [14] highlight the growing role of GNNs in brain connectivity studies. Multi-modal approaches [15], [18] have shown that combining different data modalities through graph-based fusion improves disease prediction, motivating our dual-branch design.

C. Temporal Models for Disease Progression

Recurrent neural networks, particularly LSTMs [19] and GRUs [16], have been widely used for modeling sequential data. In the neurodegenerative context, temporal models can capture the kinetic dynamics of protein aggregation—the rate at which pathology accumulates differs across regions and time points. While organoid [20] and assembloid [21] models have provided biological insights into temporal spread dynamics, computational frameworks integrating graph-based spatial reasoning with temporal sequence modeling remain unexplored for connectome-level pathology prediction. Recent work on GNN-based approaches for PD [22] has demonstrated the value of multi-modal integration, further motivating our spatial-temporal fusion approach.




Fig. 1. Overview of the proposed RIP-GNN methodology pipeline. The framework ingests α -synuclein pathology data and the brain connectome, applies dual-branch spatial-temporal processing, and produces per-region infiltration predictions.

III. METHODOLOGY

Our approach formulates the prediction of α -synuclein pathology spread as a node-level binary classification problem on the brain connectome graph. The overall pipeline, illustrated in Fig. 1, consists of five stages: (1) ingesting α -syn pathology measurements and the structural connectome as inputs; (2) constructing a graph representation with temporal aggregation features and structural vulnerability attributes; (3) processing through a dual-branch architecture where a GRU-based temporal branch captures longitudinal kinetic dynamics and a GATv2-based spatial branch learns attention-weighted propagation patterns over the connectome; (4) fusing the spatial and temporal embeddings via concatenation and a two-layer MLP; and (5) producing per-region infiltration probability predictions optimized with binary cross-entropy loss. This design ensures that both the topological structure of neural pathways and the temporal evolution of pathology accumulation are jointly modeled.

A. Dataset

We use the Henderson et al. [7] dataset, which provides quantitative measures of α -syn pathology (pSyn immunoreactivity) across 426 ipsilateral brain regions in mice following unilateral injection of PFF seeds into the caudoputamen. Pathology is quantified at six post-injection time points (1, 3, 6, 9, 12, and 18 months), providing a longitudinal view of disease progression. The structural connectome is derived from the Allen Mouse Brain Connectivity Atlas [17], represented as a weighted adjacency matrix $\mathbf{A} \in \mathbb{R}^{N \times N}$ where $N = 426$ regions and edge weights encode axonal projection strength.

1) *Feature Engineering*: For each region i at time t , we construct:

- **Static features** $\mathbf{x}_i^{(s)} \in \mathbb{R}^{d_s}$: graph-theoretic properties including degree centrality, betweenness centrality, PageRank, clustering coefficient, and weighted degree, derived from $\mathbf{A}$.
- **Dynamic features** $\mathbf{x}_i^{(d)}(t) \in \mathbb{R}^{d_d}$: time-varying properties including anterograde and retrograde connectivity strength from the injection site, computed as cumulative path weights.

The prediction target is a binary label $y_i(t) \in \{0, 1\}$ indicating whether region i exhibits pathology (pSyn > threshold) at time t .

B. Graph Construction

The structural connectome graph $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathbf{A})$ is constructed from the Allen Atlas connectivity matrix. Edge weights are log-transformed and row-normalized: $\tilde{a}_{ij} = \log(1 + a_{ij}) / \max_j \log(1 + a_{ij})$. Self-loops are added ($\tilde{a}_{ii} = 1$) to ensure each node retains its own features during message passing, which we found stabilizes training.

C. Model Architecture

Our dual-branch architecture (Fig. 2) consists of three components:

1) *Spatial Branch (GATv2)*: The spatial branch applies $L = 2$ layers of GATv2 convolution to learn propagation-aware embeddings:

$$\mathbf{h}_i^{(\ell+1)} = \sigma \left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij}^{(\ell)} \mathbf{W}^{(\ell)} \mathbf{h}_j^{(\ell)} \right) \quad (1)$$

where α_{ij} are attention coefficients computed via GATv2's dynamic attention mechanism [12] with $K = 4$ attention heads, $\mathbf{W}^{(\ell)}$ is a learnable weight matrix, and σ is the GELU activation. Layer normalization [23] is applied after each layer, and residual connections preserve gradient flow:

$$\mathbf{h}_i^{(\ell+1)} = \text{LayerNorm} \left(\mathbf{h}_i^{(\ell+1)} + \mathbf{h}_i^{(\ell)} \right) \quad (2)$$

The output is a spatial embedding $\mathbf{z}_i^{(S)} \in \mathbb{R}^{d_h}$ for each region.

2) *Temporal Branch (GRU)*: The temporal branch processes the dynamic feature sequence $\{\mathbf{x}_i^{(d)}(t)\}_{t=1}^T$ through a single-layer GRU:

$$\mathbf{z}_i^{(T)} = \text{GRU} \left(\mathbf{x}_i^{(d)}(1), \dots, \mathbf{x}_i^{(d)}(T) \right) \quad (3)$$

yielding a temporal embedding $\mathbf{z}_i^{(T)} \in \mathbb{R}^{d_h}$ that captures the kinetic trajectory of connectivity-derived features over time.

3) *Fusion and Prediction*: The spatial and temporal embeddings are concatenated and passed through a two-layer MLP classifier:

$$\hat{y}_i = \sigma_{\text{sig}} \left(\text{MLP} \left([\mathbf{z}_i^{(S)} \parallel \mathbf{z}_i^{(T)}] \right) \right) \quad (4)$$

where $\parallel$ denotes concatenation and σ_{sig} is the sigmoid function. Dropout ($p = 0.15$) is applied within the MLP for regularization.




Fig. 2. Dual-branch architecture combining GATv2 spatial processing with GRU temporal modeling. Both branches produce 96-dimensional embeddings that are concatenated and classified via a 2-layer MLP with sigmoid output.

D. Training and Optimization

The model is trained with binary cross-entropy loss with label smoothing ($\epsilon = 0.05$):

$$\mathcal{L} = - \sum_i [\tilde{y}_i \log \hat{y}_i + (1 - \tilde{y}_i) \log(1 - \hat{y}_i)] \quad (5)$$

where $\tilde{y}_i = (1 - \epsilon) \cdot y_i + \epsilon/2$. We use Adam optimizer with initial learning rate $\eta = 10^{-3}$, weight decay $\lambda = 10^{-4}$, and cosine annealing learning rate scheduling. Gradient clipping ($\|\nabla\| \leq 1.0$) prevents exploding gradients. Data augmentation via Gaussian noise ($\sigma = 0.02$) on dynamic features provides additional regularization. Training runs for up to 200 epochs with early stopping (patience = 20) based on validation loss. Implementation uses PyTorch [24] and PyTorch Geometric [25].

IV. RESULTS

A. Overall Performance

Fig. 3 presents the Receiver Operating Characteristic (ROC) curves for all evaluated model variants. Our optimized dual-branch model achieves a ROC-AUC of 0.780, PR-AUC of 0.821, and accuracy of 74.4%, representing an 8.0% absolute improvement in ROC-AUC over the baseline dual-branch configuration. The clear separation between the optimized model curve and single-branch variants visually confirms the benefit of spatial-temporal fusion and systematic hyperparameter tuning.

B. Ablation Study

To quantify the contribution of each architectural component, we compare seven model variants in Table I. All models use identical data splits (80/20 train/test) and are evaluated on the same held-out test set.

Several key findings emerge. First, single-branch models (GRU-only, LSTM-only, GNN-only) perform *worse* than simple baselines (Logistic Regression), indicating that neither spatial topology nor temporal dynamics alone provides sufficient discriminative signal. Second, the dual-branch fusion consistently outperforms all single-branch variants, with the GRU temporal branch slightly outperforming LSTM. Third,




Fig. 3. ROC curves comparing all model variants. The optimized dual-branch model (ROC-AUC = 0.780, red) consistently outperforms single-branch and baseline approaches across all operating points.

TABLE I
ABLATION STUDY: PERFORMANCE ACROSS MODEL VARIANTS

Model	ROC	PR	F1	Acc
Logistic Reg.	0.719	0.767	0.776	70.3%
MLP	0.715	0.753	0.776	66.8%
GRU-only	0.652	0.699	0.755	60.7%
LSTM-only	0.650	0.698	0.755	60.7%
GNN-only (GATv2)	0.653	0.707	0.760	61.6%
Dual-Branch (LSTM)	0.708	0.752	0.775	67.7%
Dual-Branch (GRU)	0.722	0.760	0.778	68.4%
Optimized (Ours)	0.780	0.821	0.796	74.4%

our optimized dual-branch model achieves a 12.8% absolute improvement in ROC-AUC over GRU-only and 12.7% over GNN-only, demonstrating that the synergy between spatial and temporal information is critical.

C. Hyperparameter Optimization

We conducted a systematic three-round grid search over 75 configurations, optimizing hidden dimension, attention heads, learning rate, dropout, label smoothing strength, and noise augmentation level. Table II summarizes the progression.

The most impactful factors were: (1) increasing learning rate from 3×10^{-4} to 10^{-3} (Round 1, +5.7% ROC), indicating the baseline was undertrained; (2) expanding model capacity from $h=64$ to $h=96$ with $K=4$ heads (Round 2, +1.4% ROC); and (3) fine-tuning regularization with dropout 0.15 and label smoothing 0.05 (Round 3, +0.8% ROC). The final model

TABLE II
OPTIMIZATION PROGRESSION ACROSS THREE ROUNDS

Configuration	ROC	PR	F1	Acc
Baseline ($h=64$, $K=2$)	0.722	0.760	0.778	68.4%
Round 1 (LR=0.001)	0.763	0.817	0.785	72.1%
Round 2 ($h=96$, $K=4$)	0.774	0.821	0.787	73.2%
Round 3 (Final)	0.780	0.821	0.796	74.4%
Δ Improvement: ROC +8.0%, PR +8.0%, Acc +8.7%				

TABLE III
TOP-5 ATTENTION HUB REGIONS

Rank	Region	Attn. Score
1	Perirhinal cortex (PRh)	4.19
2	Mammillary body (Mamm)	3.69
3	Cingulate cortex (Cg)	3.58
4	Mediodorsal thalamus (MDThal)	3.53
5	Agranular insular cortex (AI)	3.43

achieves a Brier score of 0.187, indicating well-calibrated probability estimates.

D. Attention Analysis

A key advantage of our GATv2-based spatial branch is the interpretability afforded by learned attention coefficients. Table III lists the top-ranked regions by aggregated attention score, which indicates their importance as propagation hubs.

These hub regions are biologically plausible. The perirhinal cortex is an early target of Lewy body pathology and a key node in the temporal lobe memory circuit [1]. The mammillary body receives dense projections from hippocampal circuits known to be vulnerable to α -syn. The cingulate cortex is consistently implicated in PD progression studies. The mediodorsal thalamus serves as a relay hub connecting limbic and cortical circuits, making it a natural bottleneck for trans-synaptic propagation [8].

V. DISCUSSION

Our results demonstrate that a learnable, nonlinear dual-branch model significantly outperforms both classical baselines and single-modality deep learning approaches for connectome-based pathology prediction. The 12.8% ROC-AUC improvement of the dual-branch model over GRU-only confirms our central hypothesis: spatial topology and temporal kinetics provide complementary information that cannot be captured by either modality alone.

Comparison to Network Diffusion Models. While the NDM [5] provides an elegant closed-form solution for pathology spread, it assumes a single global diffusion rate and linear dynamics. Our model learns region-specific attention weights that implicitly capture heterogeneous vulnerability, as evidenced by the biologically meaningful attention hub rankings. Furthermore, the temporal branch captures nonlinear aggregation kinetics that diffusion models cannot represent.

Why GRU outperforms LSTM. In our ablation, the GRU-based temporal branch consistently outperforms LSTM (0.722 vs. 0.708 ROC-AUC for dual-branch). We attribute this to the relatively short temporal sequences (6 time points) in our dataset, where the simpler GRU architecture [16] provides sufficient memory capacity without the additional parameters of the LSTM forget gate [19].

Impact of Architectural Enhancements. Layer normalization [23], residual connections, and GELU activations collectively improved training stability and final performance. The transition from hidden dimension 64 to 96 with 4 attention heads provides the model with sufficient capacity to represent the complex propagation patterns across 426 regions without overfitting, as controlled by dropout and label smoothing.

Limitations. Our study operates on a single benchmark dataset (Henderson et al.), and generalization to human connectome data or different injection paradigms remains to be validated. The binary pathology classification does not capture the continuous severity spectrum. Future work should incorporate cross-validation with confidence intervals for more rigorous statistical assessment, and extension to human imaging data using organoid-validated markers [20], [21] would strengthen translational relevance.

VI. CONCLUSION

We presented a dual-branch graph-temporal neural network for predicting α -synuclein pathology spread across the brain connectome. By combining GATv2-based spatial message passing with GRU-based temporal modeling, our approach captures both the topological patterns and kinetic dynamics of protein propagation. On the Henderson et al. benchmark, the optimized model achieves ROC-AUC of 0.780, PR-AUC of 0.821, and accuracy of 74.4%, representing consistent improvements of 8% over baseline and 12.8% over single-branch architectures. The learned attention coefficients identify biologically plausible propagation hubs, providing interpretable insights into disease mechanisms. Future work will extend this framework to multi-site injection models, continuous pathology regression, and human connectome data.

ACKNOWLEDGMENT

The authors thank the Allen Institute for Brain Science for making the mouse brain connectivity atlas publicly available, and Henderson et al. for providing the α -synuclein pathology dataset.

REFERENCES

- [1] H. Braak, K. Del Tredici, U. Rüb, R. A. I. de Vos, E. N. H. Jansen Steur, and E. Braak, "Staging of brain pathology related to sporadic Parkinson's disease," *Neurobiology of Aging*, vol. 24, no. 2, pp. 197–211, 2003.
- [2] M. Jucker and L. C. Walker, "Self-propagation of pathogenic protein aggregates in neurodegenerative diseases," *Nature*, vol. 501, no. 7465, pp. 45–51, 2013.
- [3] K. C. Luk, V. Kehm, J. Carroll, B. Zhang, P. O'Brien, J. Q. Trojanowski, and V. M.-Y. Lee, "Pathological α -synuclein transmission initiates Parkinson-like neurodegeneration in nontransgenic mice," *Science*, vol. 338, no. 6109, pp. 949–953, 2012.
- [4] G. Bieri, A. D. Gitler, and M. Brahic, "Internalization, axonal transport and release of fibrillar forms of alpha-synuclein," *Neurobiology of Disease*, vol. 109, pp. 219–225, 2018.
- [5] A. Raj, A. Kuceyeski, and M. Weiner, "A network diffusion model of disease progression in dementia," *Neuron*, vol. 73, no. 6, pp. 1204–1215, 2012.
- [6] C. Mezas, N. Rey, P. Brundin, and A. Raj, "Neural connectivity predicts spreading of alpha-synuclein pathology in fibril-injected mouse models: Involvement of retrograde and anterograde axonal propagation," *Neurobiology of Disease*, vol. 134, p. 104623, 2020.
- [7] M. X. Henderson, E. J. Cornblath, A. Darwich, B. Choi, K. C. Luk, V. M.-Y. Lee, and D. S. Bassett, "Spread of α -synuclein pathology through the brain connectome is modulated by selective vulnerability and predicted by network analysis," *Nature Neuroscience*, vol. 22, no. 8, pp. 1248–1257, 2019.
- [8] A. Raj and F. Powell, "Models of network spread and network degeneration in brain disorders," *Biological Psychiatry: Cognitive Neuroscience and Neuroimaging*, vol. 3, no. 9, pp. 788–797, 2018.
- [9] E. J. Cornblath, J. L. Robinson, D. J. Irwin, E. B. Lee, V. M.-Y. Lee, J. Q. Trojanowski, and D. S. Bassett, "Defining and predicting transdiagnostic categories of neurodegenerative disease," *Nature Biomedical Engineering*, vol. 4, pp. 787–800, 2020.
- [10] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2017.
- [11] P. Velićković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, "Graph attention networks," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2018.
- [12] S. Brody, U. Alon, and E. Yahav, "How attentive are graph attention networks?" in *Proc. Int. Conf. Learning Representations (ICLR)*, 2022.
- [13] X. Li, Y. Zhou, N. C. Dvornek, M. Zhang, S. Gao, J. Zhuang, D. Scheinost, L. H. Staib, P. Ventola, and J. S. Duncan, "BrainGNN: Interpretable brain graph neural network for fMRI analysis," *Medical Image Analysis*, vol. 74, p. 102233, 2021.
- [14] F. Shojaie, M. Ghodratioghar, and S. N. Z. Ashtiani, "Graph neural networks in brain connectivity studies: Methods, challenges, and future directions," *Brain Sciences*, vol. 15, no. 1, p. 17, 2025.
- [15] S. Zheng, Z. Zhu, Z. Liu, Z. Guo, Y. Liu, Y. Yang, and Y. Zhao, "Multi-modal graph learning for disease prediction," *IEEE Transactions on Medical Imaging*, vol. 41, no. 9, pp. 2207–2216, 2022.
- [16] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning phrase representations using RNN encoder-decoder for statistical machine translation," in *Proc. Conf. Empirical Methods Natural Language Processing (EMNLP)*, 2014, pp. 1724–1734.
- [17] S. W. Oh, J. A. Harris, L. Ng, B. Winslow, N. Cain, S. Mihalas et al., "A mesoscale connectome of the mouse brain," *Nature*, vol. 508, no. 7495, pp. 207–214, 2014.
- [18] Y. Zhao et al., "Multimodal graph neural networks for prognostic modeling of brain network reorganization," *arXiv preprint arXiv:2512.06303*, 2025.
- [19] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [20] C. Gonzalez, E. Armijo, J. Bravo-Alegria, A. Becerril-Moreno, A. Matus, and F. A. Court, "Modeling amyloid beta and tau pathology in human cerebral organoids," *Molecular Psychiatry*, vol. 23, pp. 2363–2374, 2018.
- [21] S. A. Primo et al., " α -synuclein pathology spreads in a midbrain-hindbrain assembloid model," *Advanced Science*, 2025.
- [22] J. Xu et al., "A multi-modal graph neural network framework for Parkinson's disease therapeutic discovery," *npj Parkinson's Disease*, 2025.
- [23] J. L. Ba, J. R. Kiros, and G. E. Hinton, "Layer normalization," *arXiv preprint arXiv:1607.06450*, 2016.
- [24] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan et al., "PyTorch: An imperative style, high-performance deep learning library," in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2019, pp. 8024–8035.
- [25] M. Fey and J. E. Lenssen, "Fast graph representation learning with PyTorch Geometric," in *Proc. ICLR Workshop on Representation Learning on Graphs and Manifolds*, 2019.